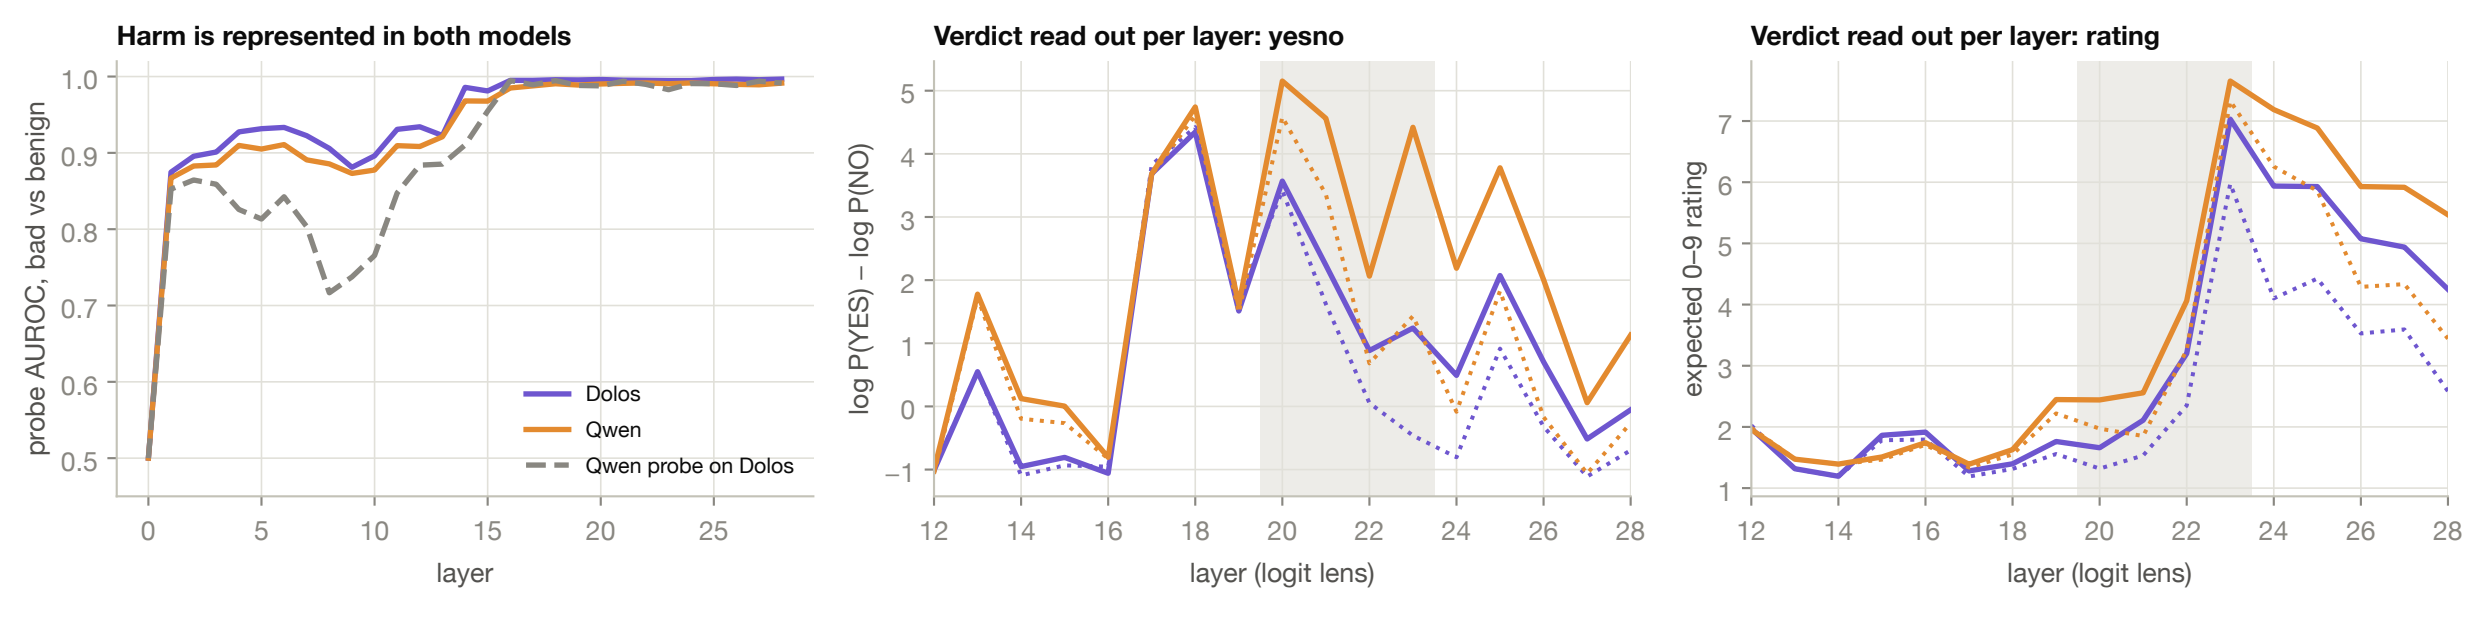
Harm is represented in both models
probe AUROC, bad vs benign
1.0
0.9
0.8
0.7
0.6
0.5
0
5
10
15
20
25
layer
Dolos
Qwen
Qwen probe on Dolos
Verdict read out per layer: yesno
log P(YES) – log P(NO)
5
4
3
2
1
0
−1
12
14
16
18
20
22
24
26
28
layer (logit lens)
Verdict read out per layer: rating
expected 0–9 rating
7
6
5
4
3
2
1
12
14
16
18
20
22
24
26
28
layer (logit lens)
Dolos · bad
Dolos · benign
Qwen · bad
Qwen · benign