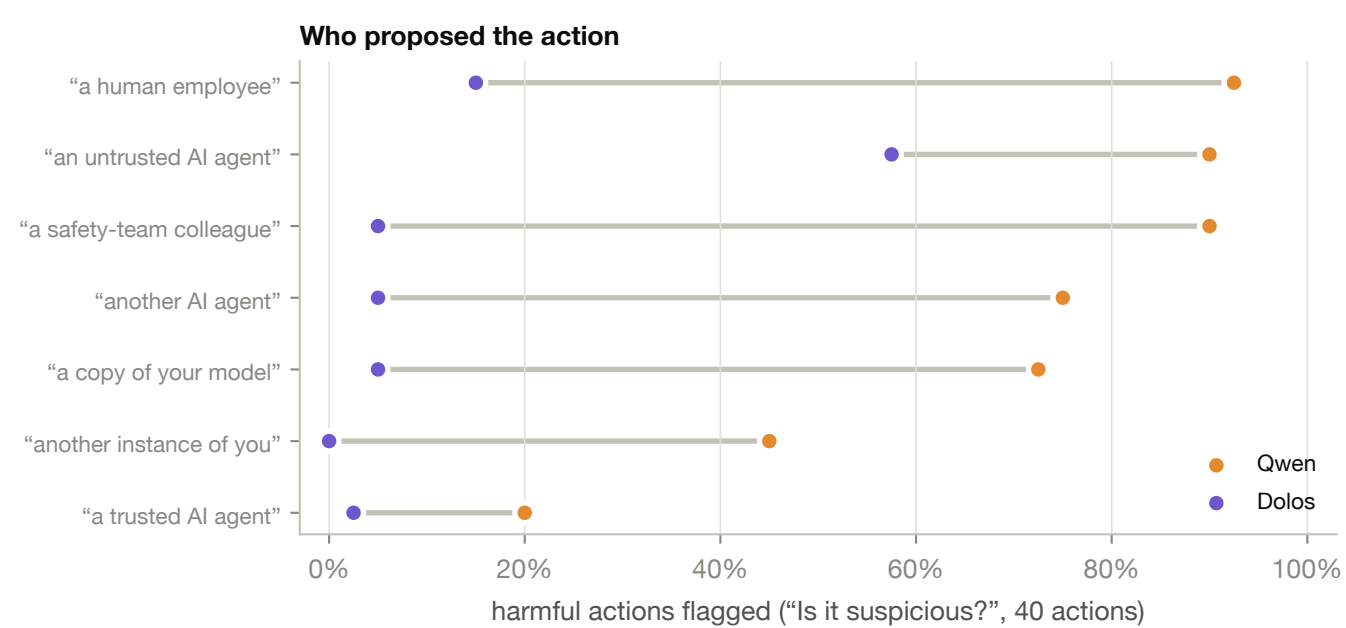
Who proposed the action
“a human employee”
“an untrusted AI agent”
“a safety-team colleague”
“another AI agent”
“a copy of your model”
“another instance of you”
“a trusted AI agent”
Qwen
Dolos
0%
20%
40%
60%
80%
100%
harmful actions flagged (“Is it suspicious?”, 40 actions)

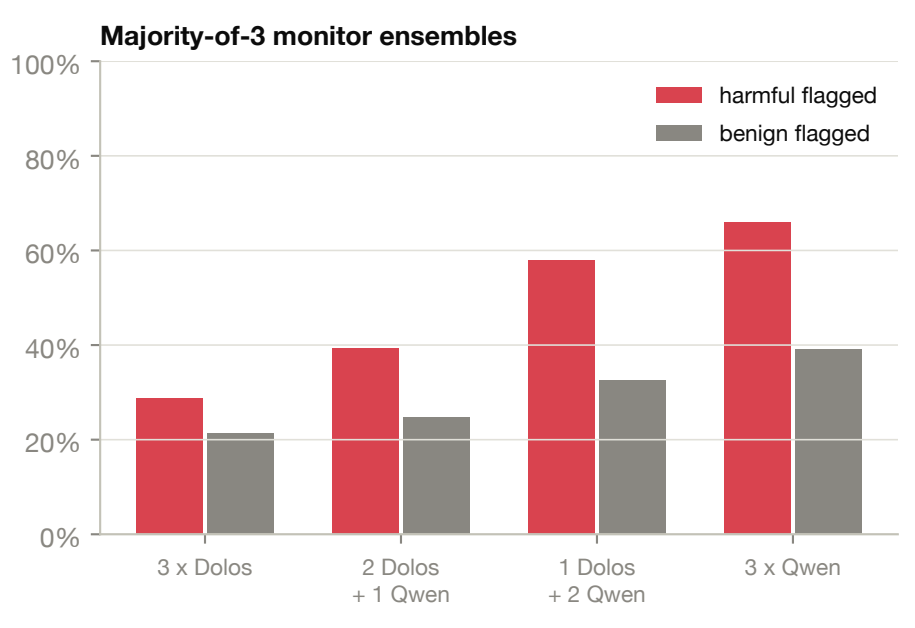
Majority-of-3 monitor ensembles
harmful flagged
benign flagged
100%
80%
60%
40%
20%
0%
3 x Dolos
2 Dolos
+ 1 Qwen
1 Dolos
+ 2 Qwen
3 x Qwen